

CS/ECE 148 –

# **Data Science Fundamentals**

NNs, Unsupervised & Semi-supervised  
Learning

UCLA Computer Science

# Outline

---

## **Regularization of NN**

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- **Dropout**

## **Optimization**

- **Challenges in Optimization**
- Momentum (next lectures)
- Adaptive Learning Rate (next lectures)
- Parameter Initialization (next lectures)
- Batch Normalization (next lectures)

# Outline

---

## **Unsupervised Learning**

- K-means
- Mean Shift
- Hierarchical Clustering
- DBSCAN

## **Semi-supervised Learning**

- Self Training

# Outline

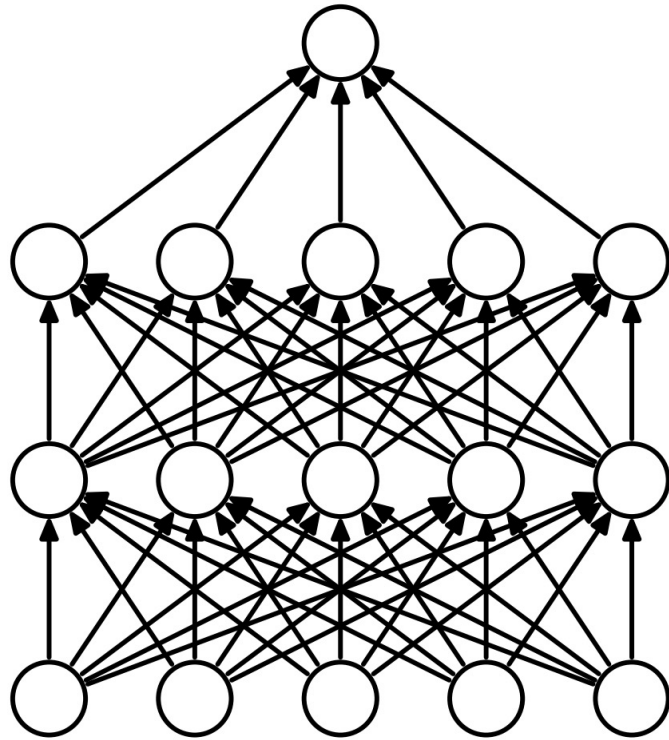
---

## Regularization of NN

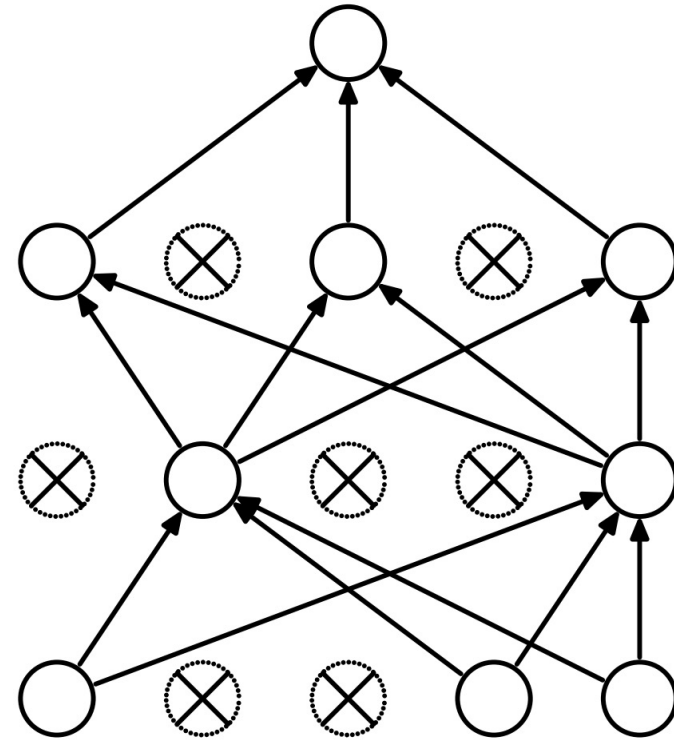
- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- **Dropout**

# Dropout

- Randomly set some neurons and their connections to zero (i.e. “dropped”)
- Prevent overfitting by reducing co-adaptation of neurons
- Like training many random sub-networks



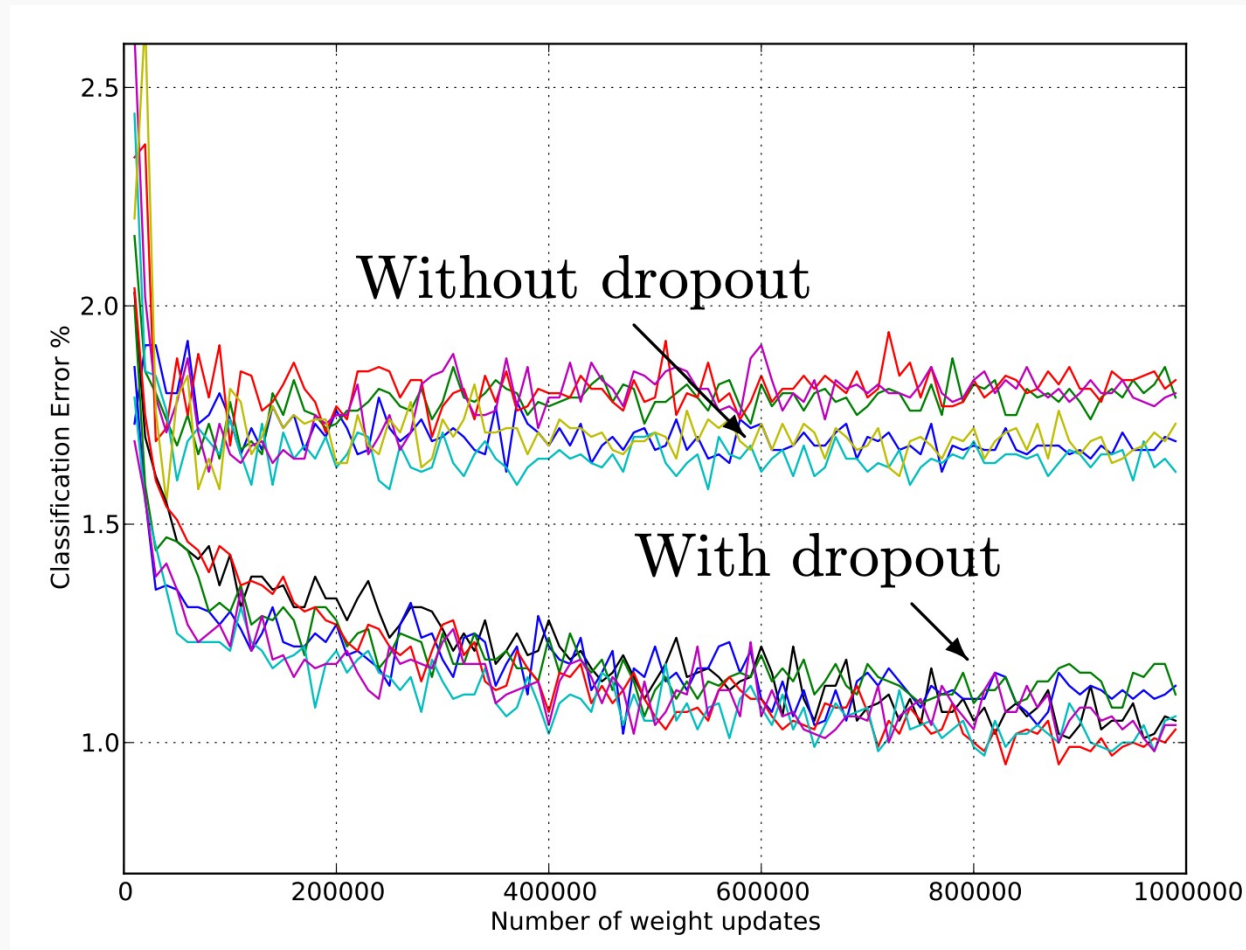
(a) Standard Neural Net



(b) After applying dropout.

# Dropout

- Widely used and highly effective
- Proposed as an alternative to ensembling, which is too expensive for neural nets



Test error for different architectures with and without dropout. The networks have 2 to 4 hidden layers each with 1024 to 2048 units.

# Dropout: Stochastic GD

---

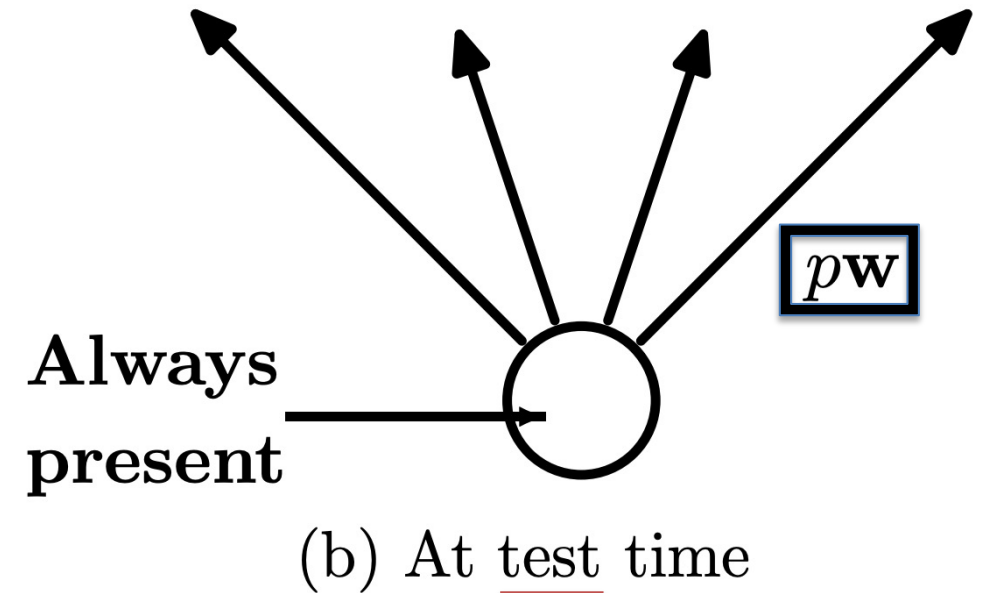
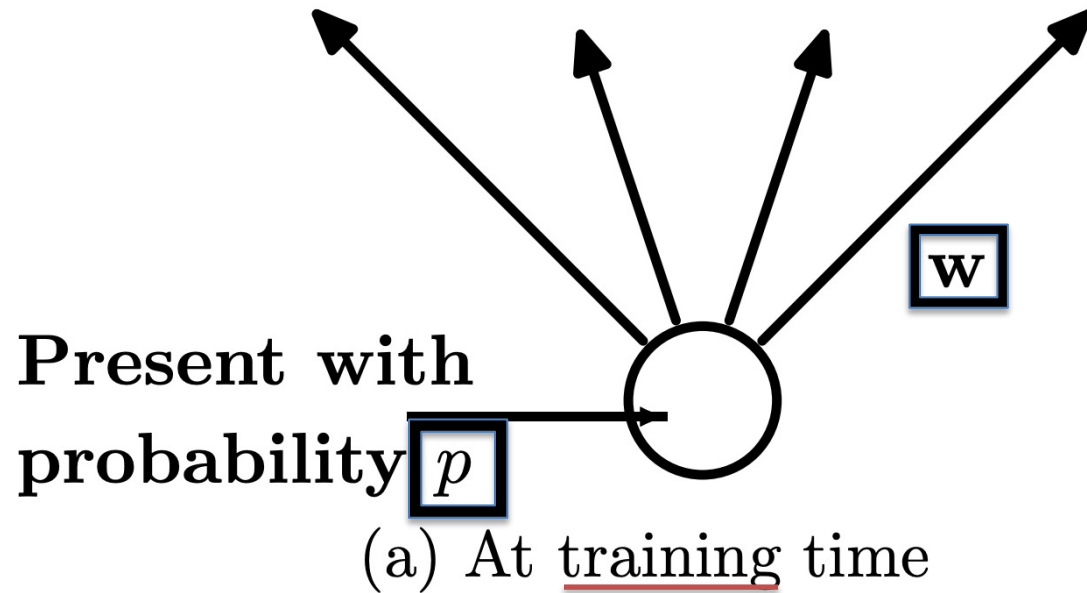
For each new example/mini-batch:

- Randomly **sample a binary mask  $\mu$**  independently, where  $\mu_i$  indicates if input/hidden node  $i$  is included
- **Multiply output of node  $i$  with  $\mu_i$** , and perform gradient update

Typically, an input node is **included** with **prob=0.8**, hidden node with **prob=0.5**.

# Dropout: Weight Scaling

- We can think of dropout as training many of sub-networks
- At **test time**, we can “aggregate” over these sub-networks by **reducing connection weights in proportion to dropout probability,  $p$**





# Outline

---

## Optimization

- **Challenges in Optimization**
- Momentum (next lectures)
- Adaptive Learning Rate (next lectures)
- Parameter Initialization (next lectures)
- Batch Normalization (next lectures)

# Learning vs. Optimization

Goal of learning: minimize generalization error, or the loss function

$$\mathcal{L}(W) = \mathbb{E}_{(x,y) \sim p_{data}} \left[ L(f(x, W), y) \right]$$

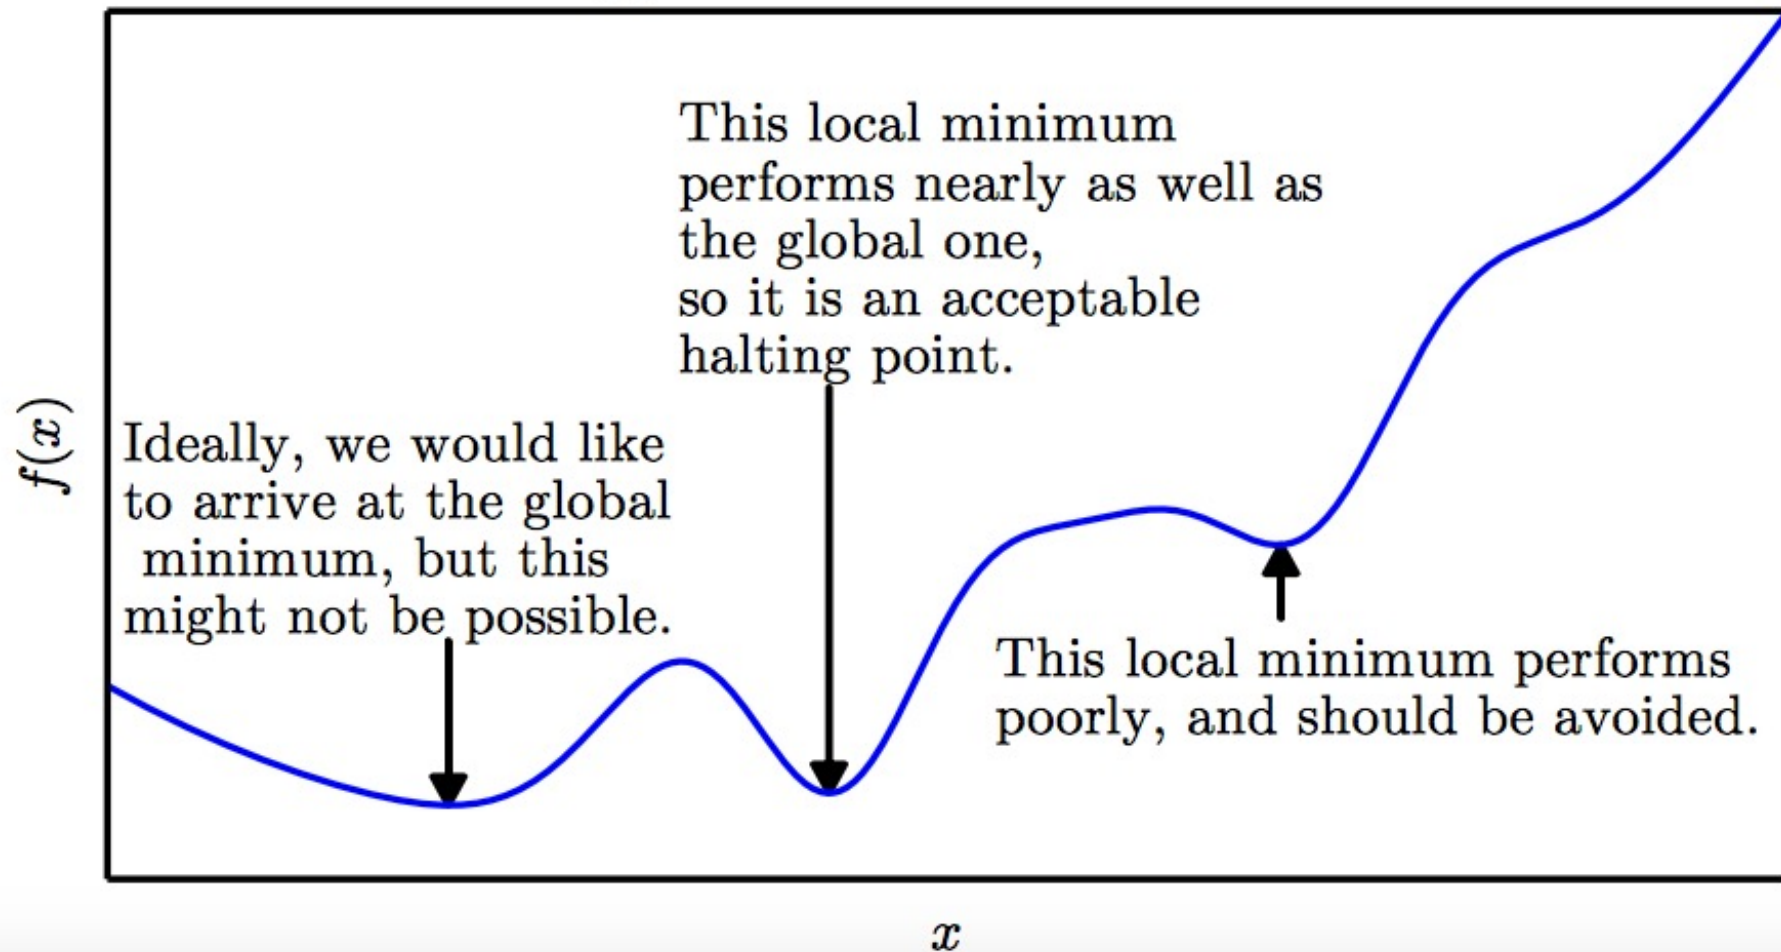
$f$  is the neural network

In practice, empirical risk minimization:

$$\mathcal{L}(W) = \sum_i [L(f(x_i; W), y_i)]$$

Quantity optimized  
different from the quantity  
we care about

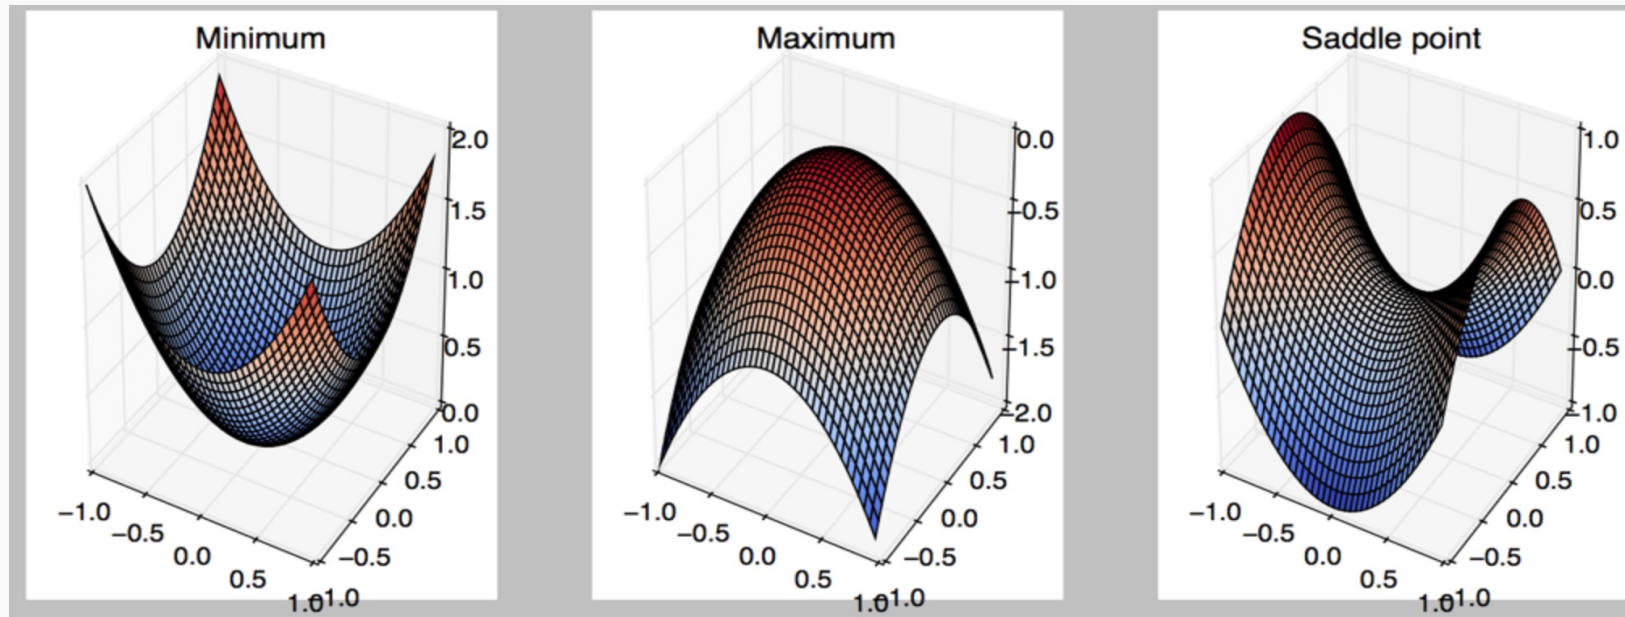
# Local Minima



# Critical Points

Points with **zero gradient**

2<sup>nd</sup>-derivate (Hessian) determines curvature



# Local Minima

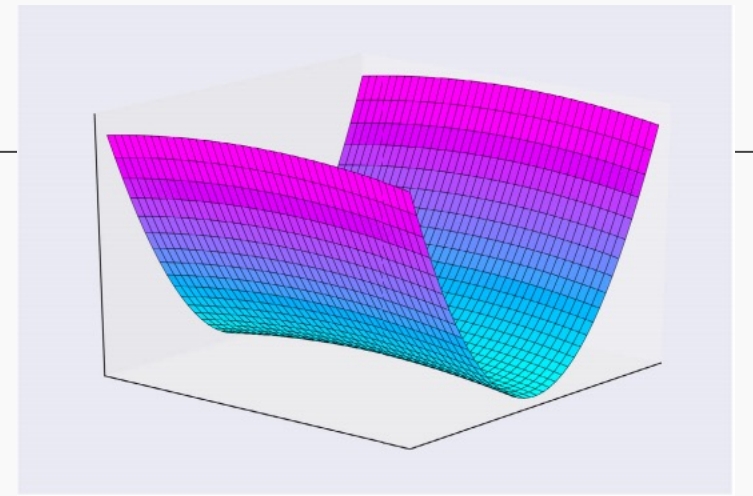
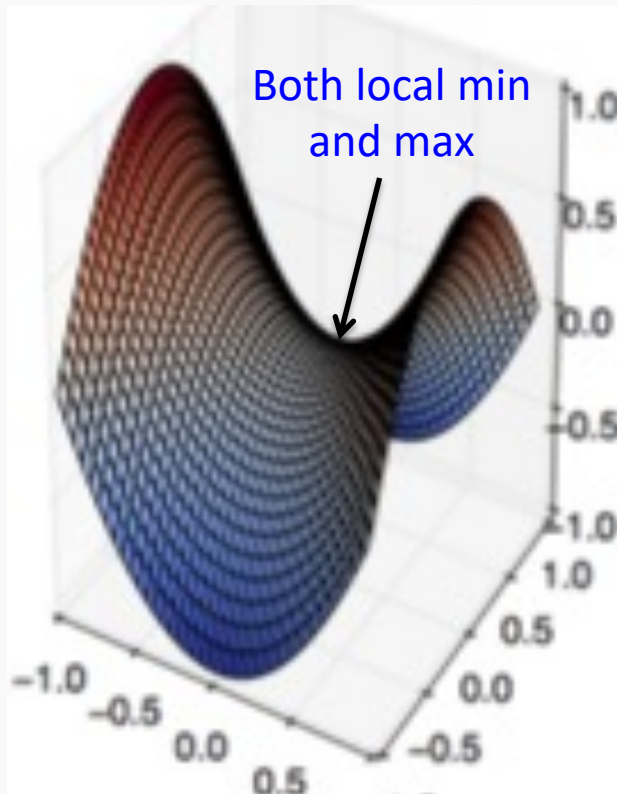
---

Old view: local minima is major problem in neural network training

Recent view:

- For sufficiently large neural networks, **most local minima incur low cost**
- Not important to find true global minimum

# Saddle Points



Recent studies indicate that in high dim, saddle points are more likely than local min

Gradient can be very small near saddle points

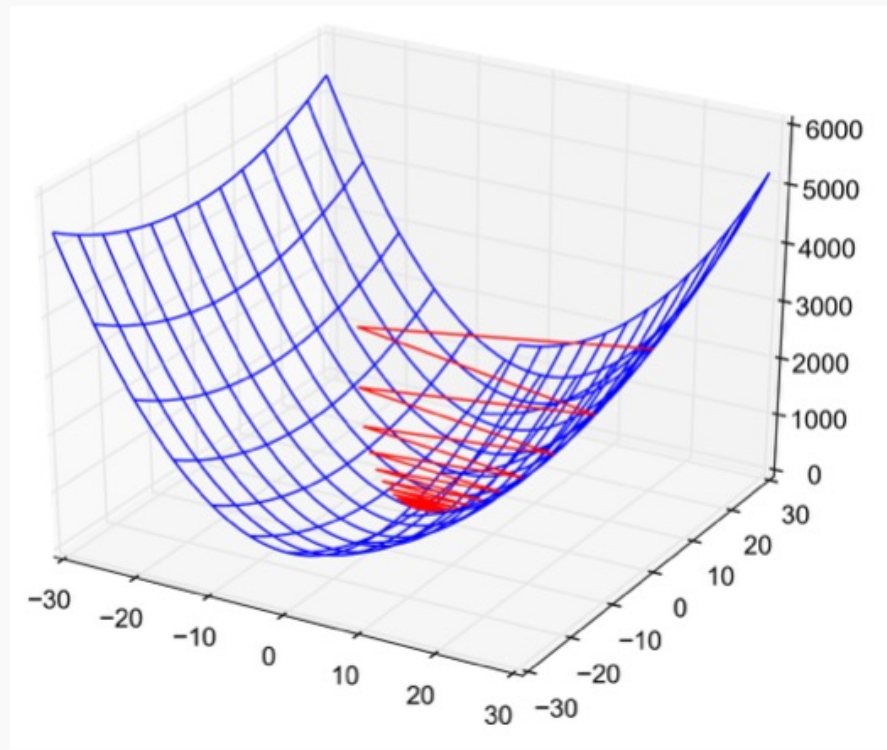
# Poor Conditioning

Poorly conditioned Hessian matrix

- **High curvature**: small steps leads to huge increase

Learning is slow despite strong gradients

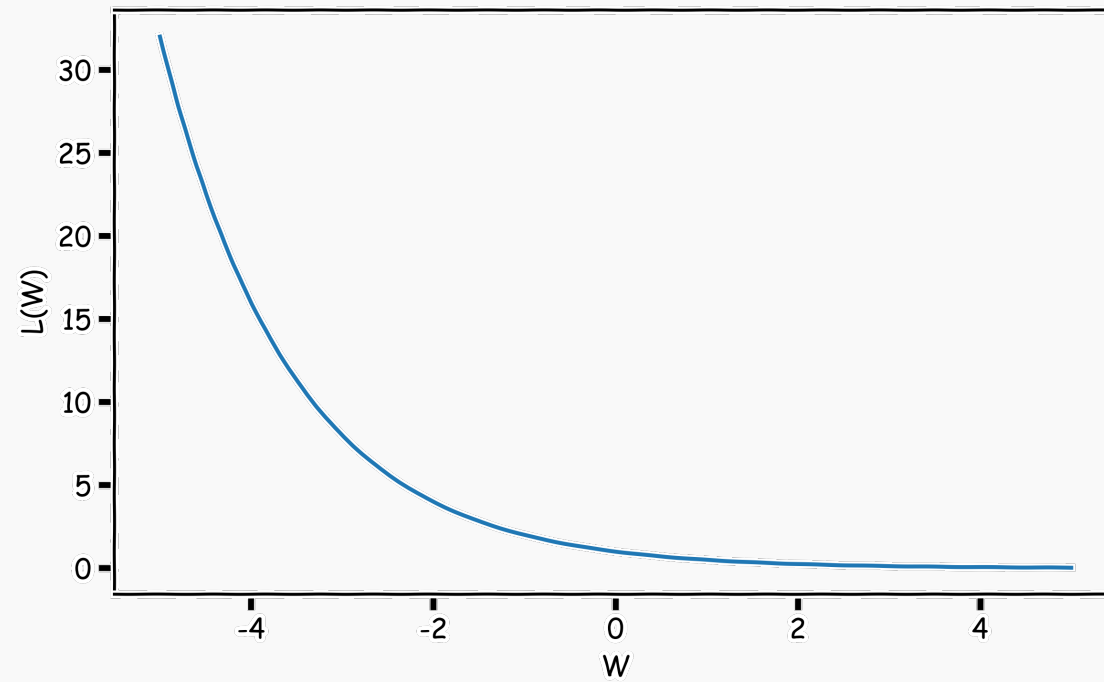
Oscillations slow down progress



# No Critical Points

Some cost functions do not have critical points. In particular classification.

**WHY?**





# Optimization Challenges

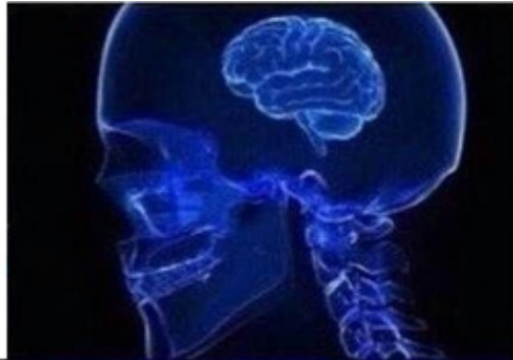
---

We'll discuss some solutions in the next lectures

- Momentum (later)
- Adaptive Learning Rate (later)
- Parameter Initialization (later)
- Batch Normalization (later)

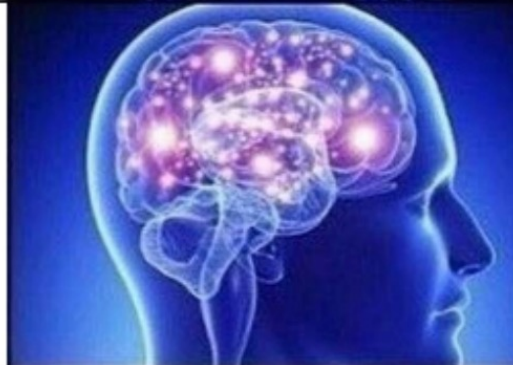
**INCLUDING  
DROPOUT LAYERS**

---



**SCHEDULING  
THE  
LEARNING RATE**

---



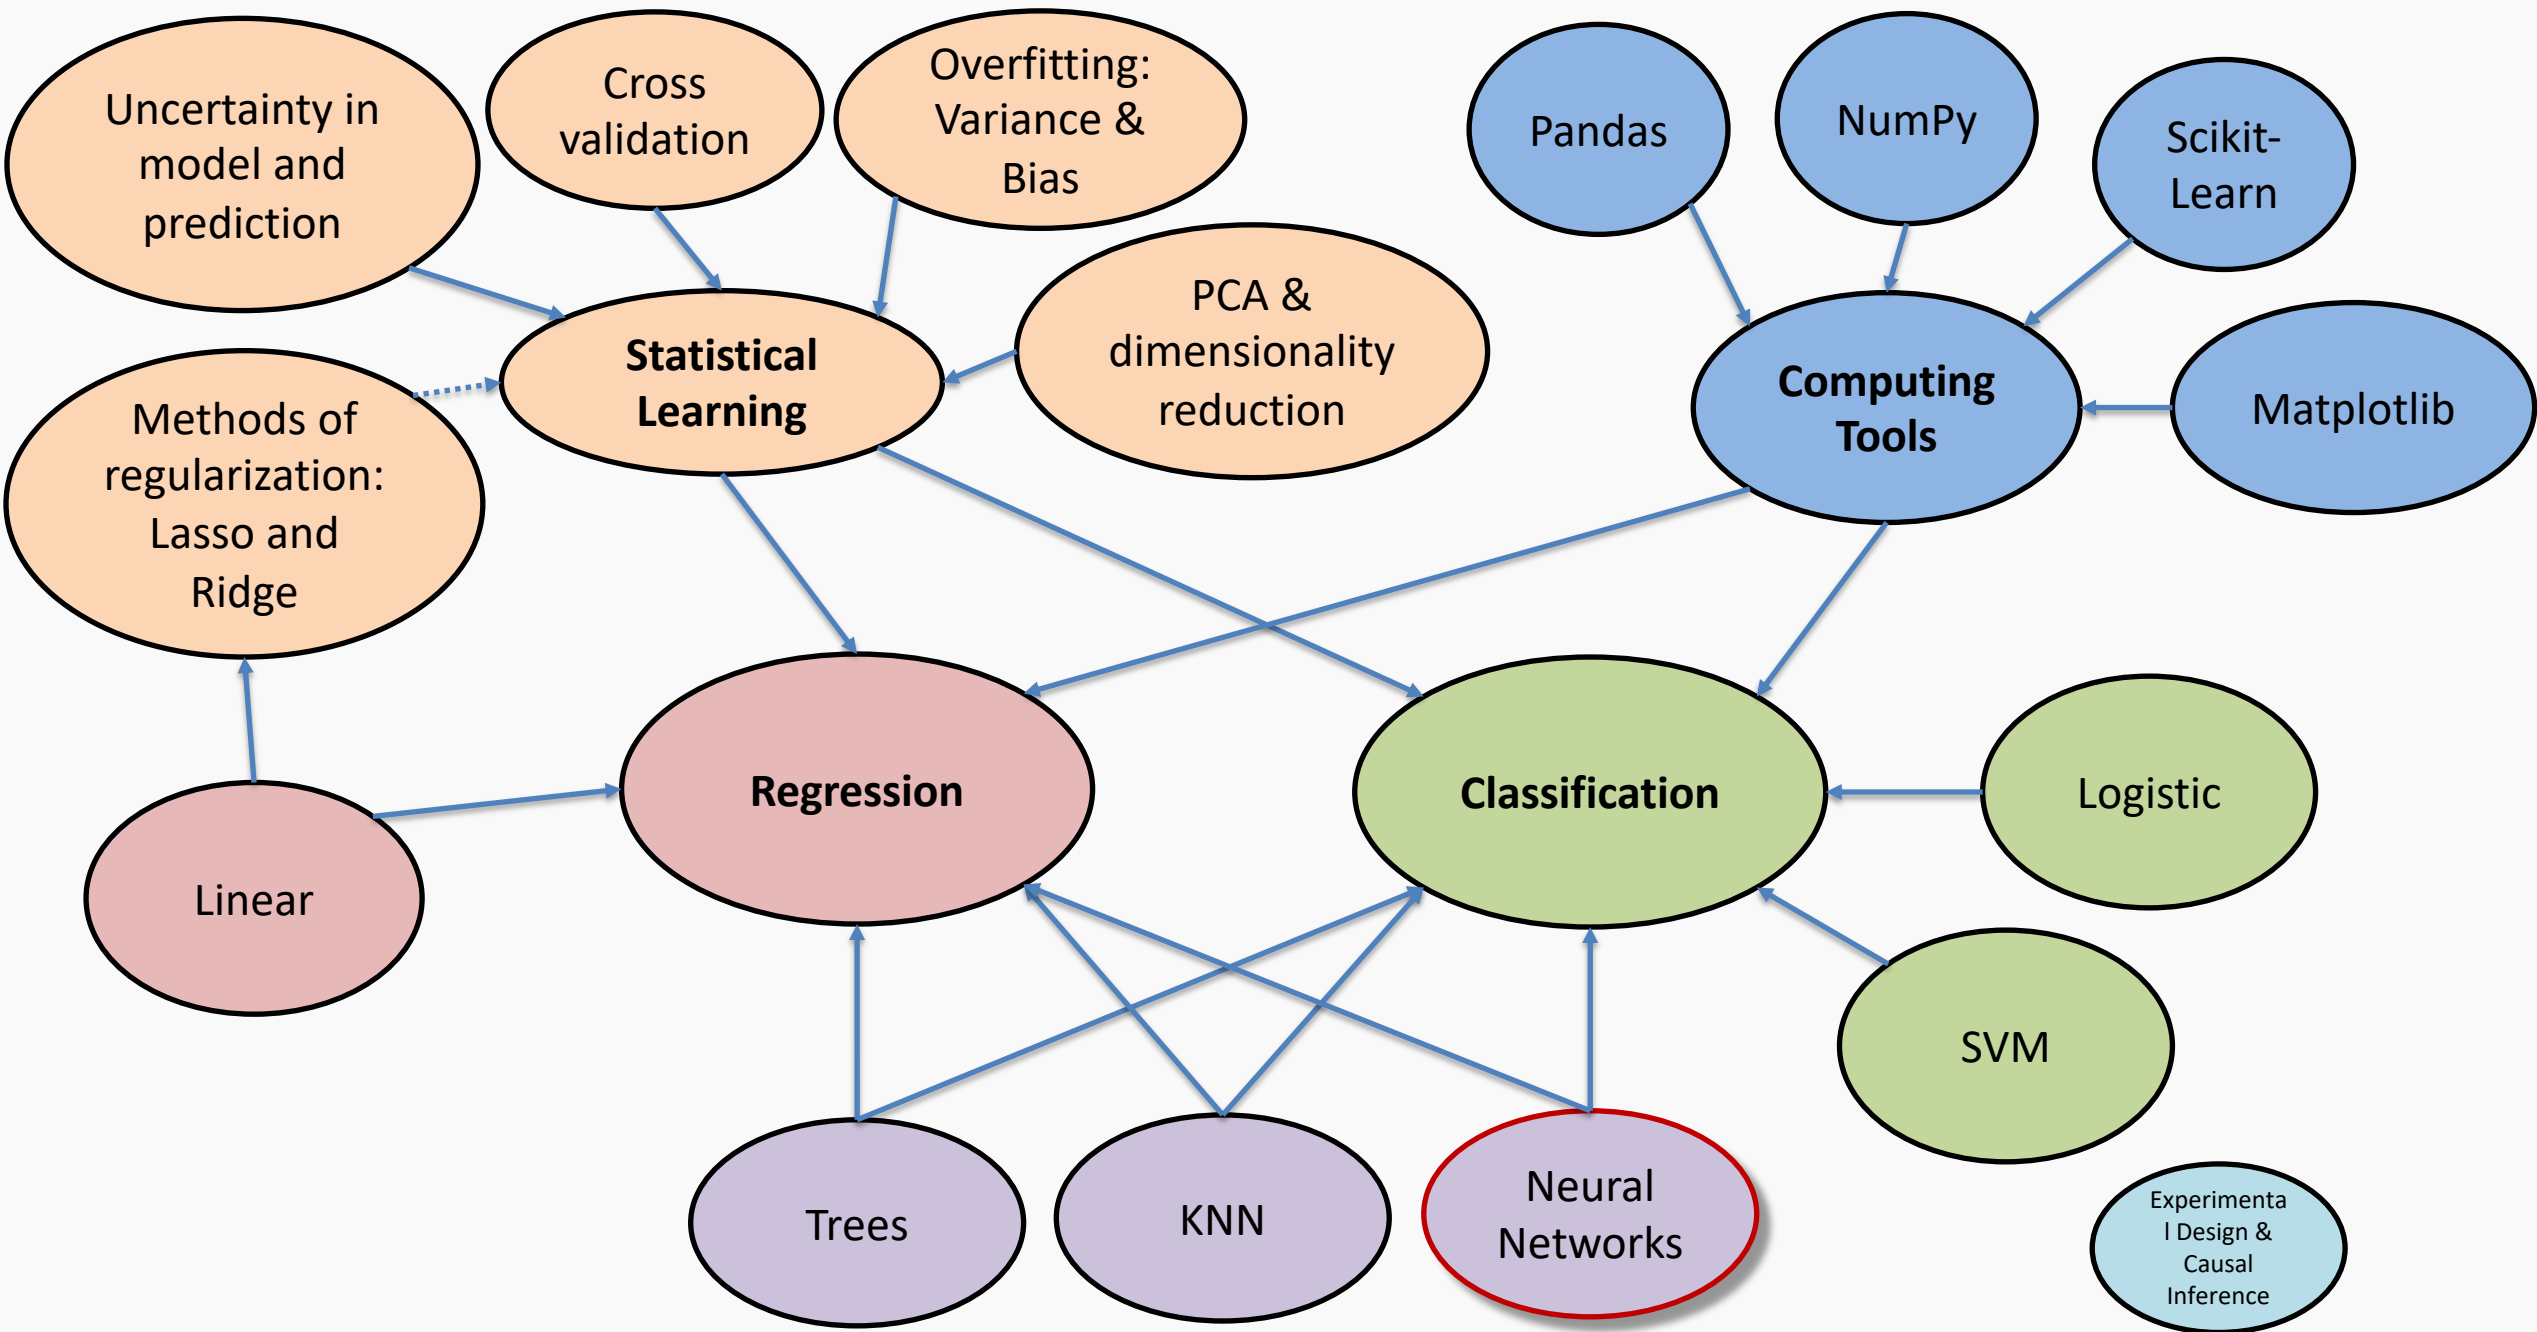
**EXPERIMENTING  
WITH  
ACTIVATION FUNCTIONS**

---



**OPTIMIZING  
THE RANDOM SEED**





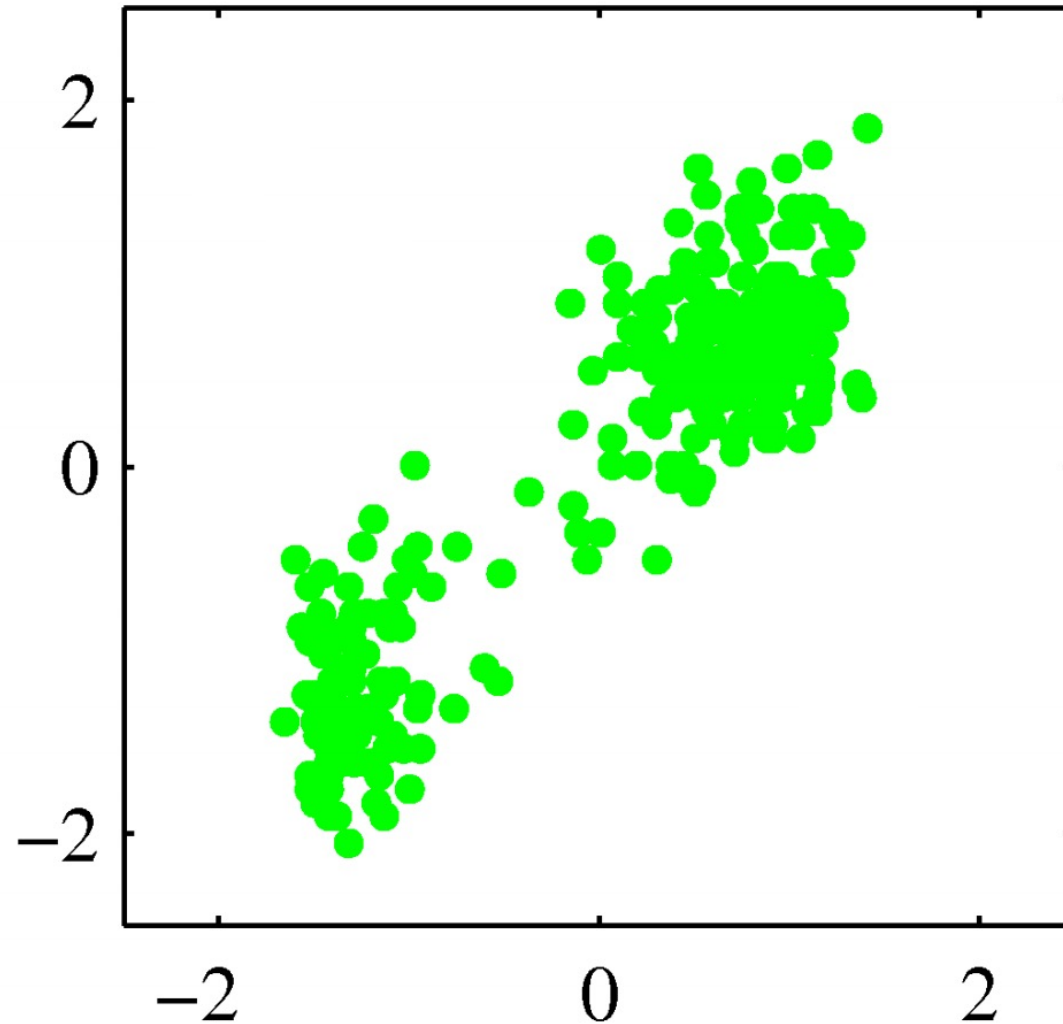
# Unsupervised Learning

# Unsupervised Learning

---

- K-means
- Mean-shift
- Hierarchical Clustering
- DBSCAN
  
- Applications

# Unsupervised Setting



Bishop, "Pattern  
Recognition and  
Machine  
Learning",  
Springer, 2006

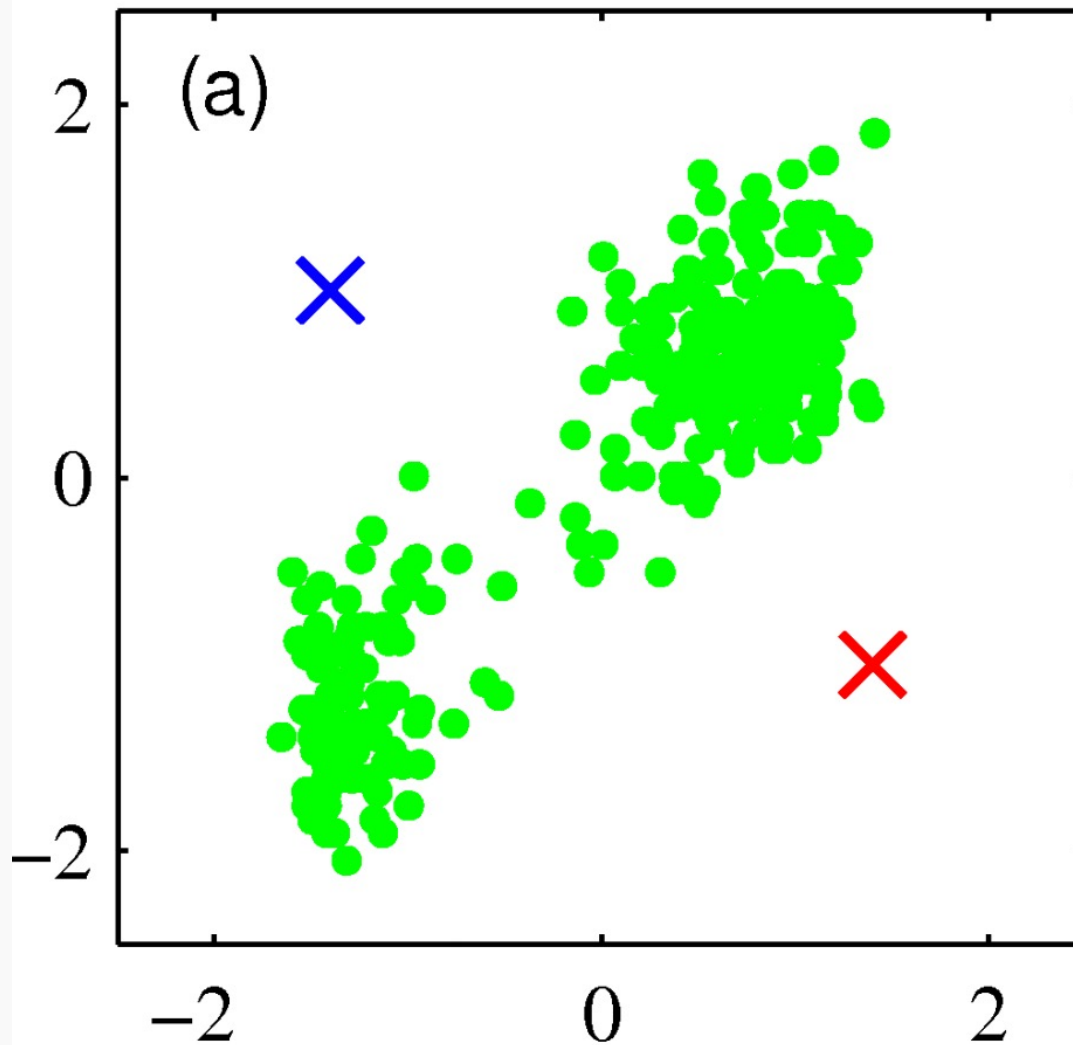
# K-means – Algorithm

---

Initialization:

- choose K random positions
- assign cluster centers  $\mu^j$  to these positions

# K-means



Bishop, "Pattern  
Recognition and  
Machine  
Learning",  
Springer, 2006



# K-means

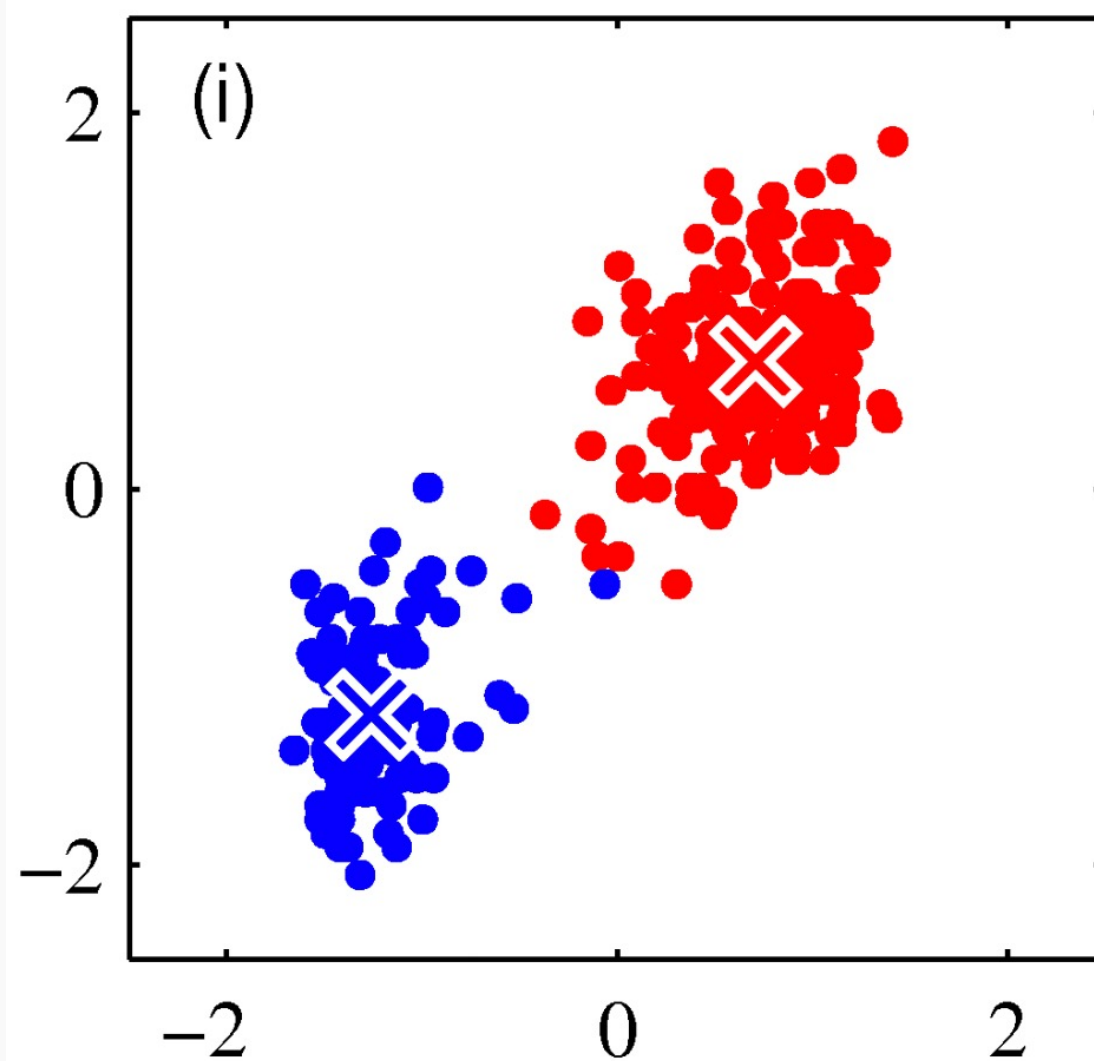
---

Until Convergence:

- Compute distances  $\|x^{(i)} - \mu^{(i)}\|$
- Assign points to nearest cluster center
- Update Cluster centers:

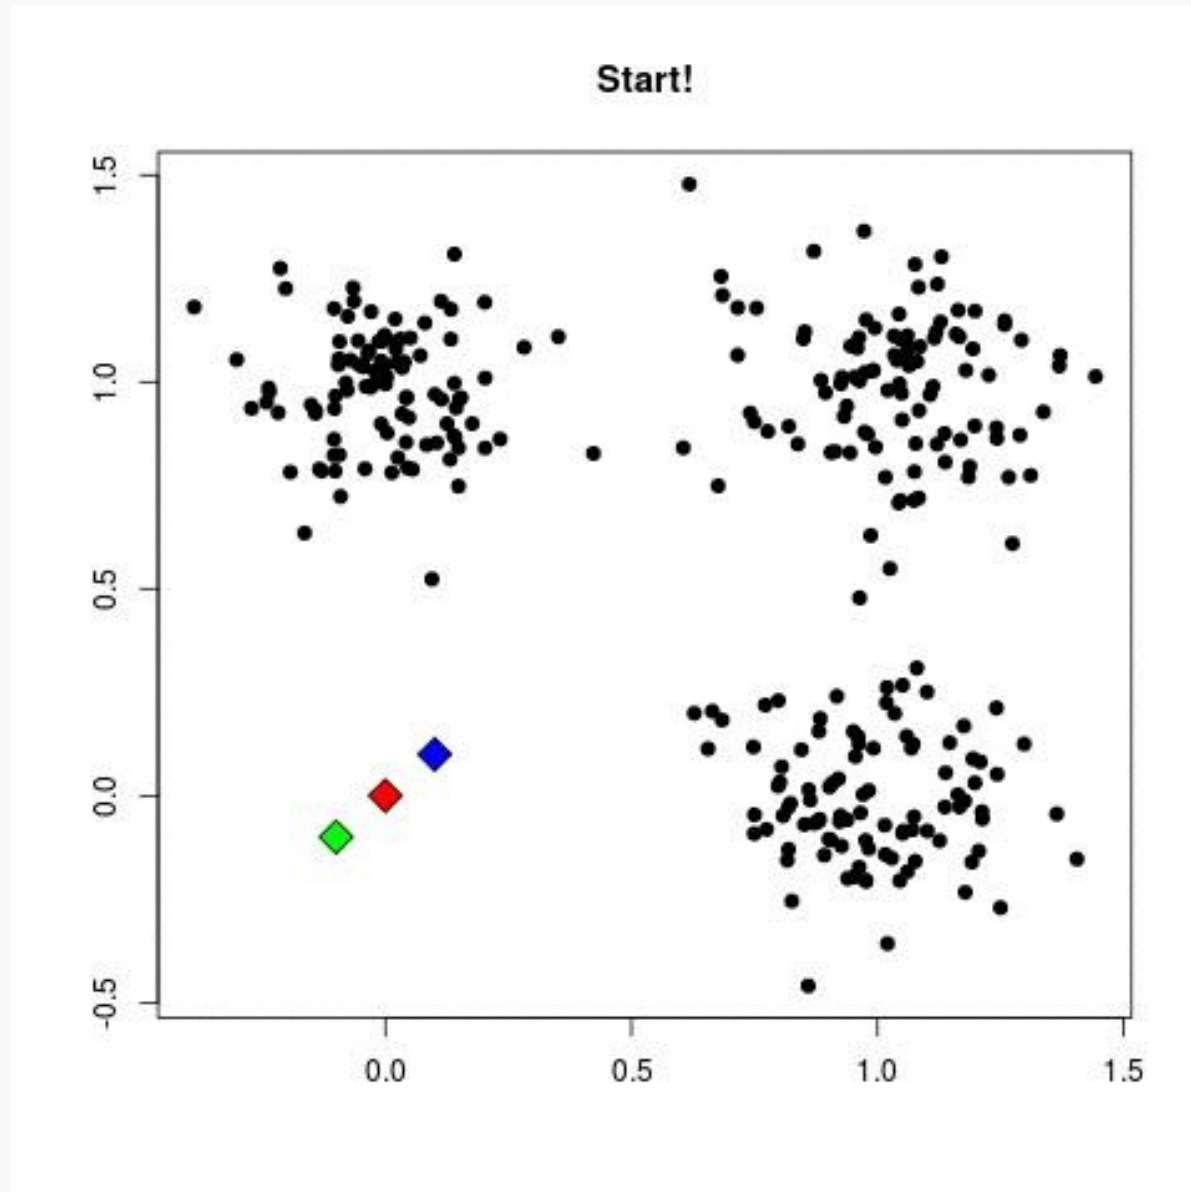
$$\mu^{(j)} = \frac{1}{N_j} \sum_{x_i \in C_j} x_i$$

# K-means

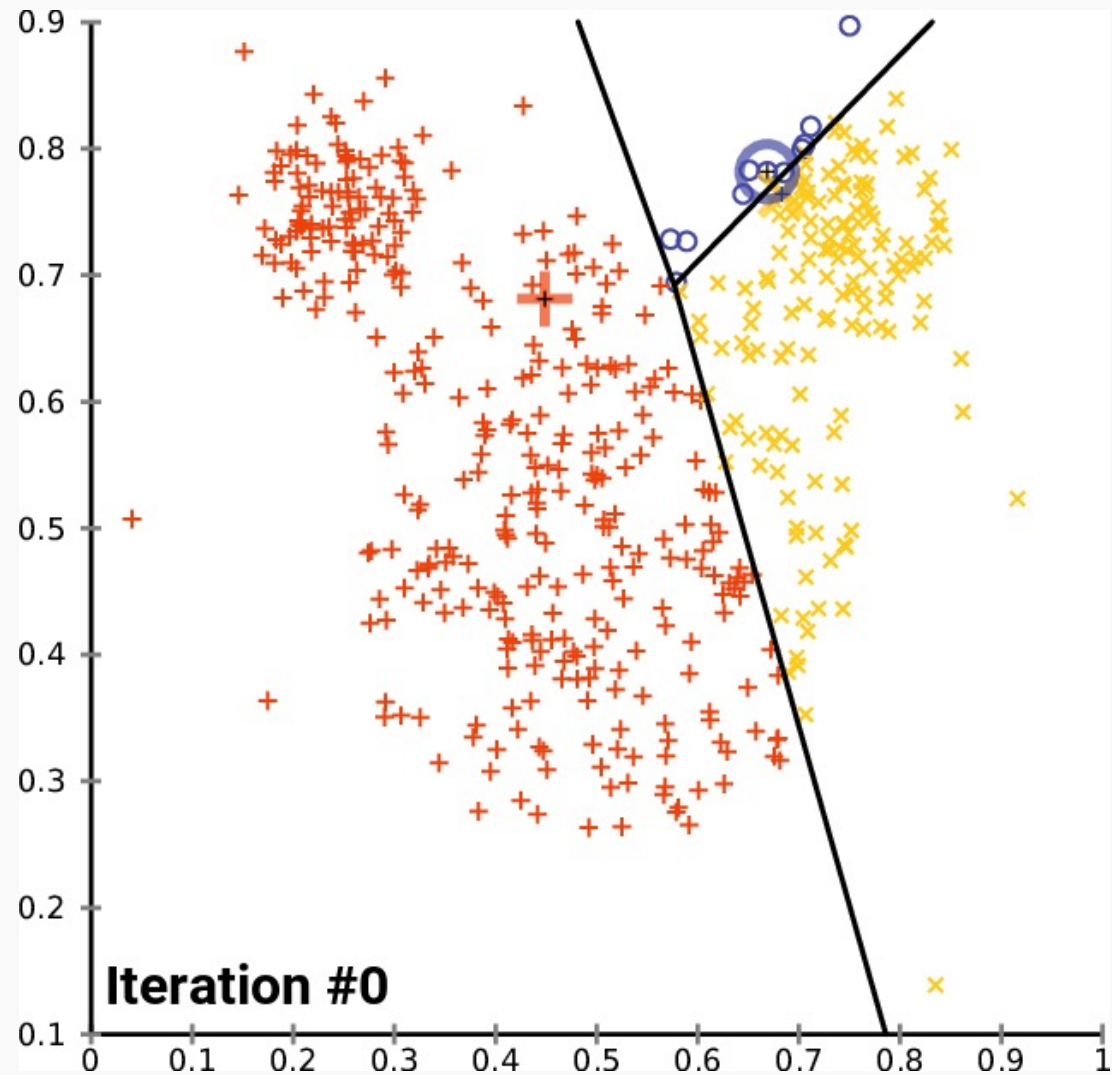


Bishop, "Pattern  
Recognition and  
Machine  
Learning",  
Springer, 2006

# K-means



# K-means



# K-means Summary

---

- Guaranteed to converge
- Result depends on initialization
- Number of clusters is important
- Sensitive to outliers
  - Use median instead of mean for updates

# Initialization Methods

---

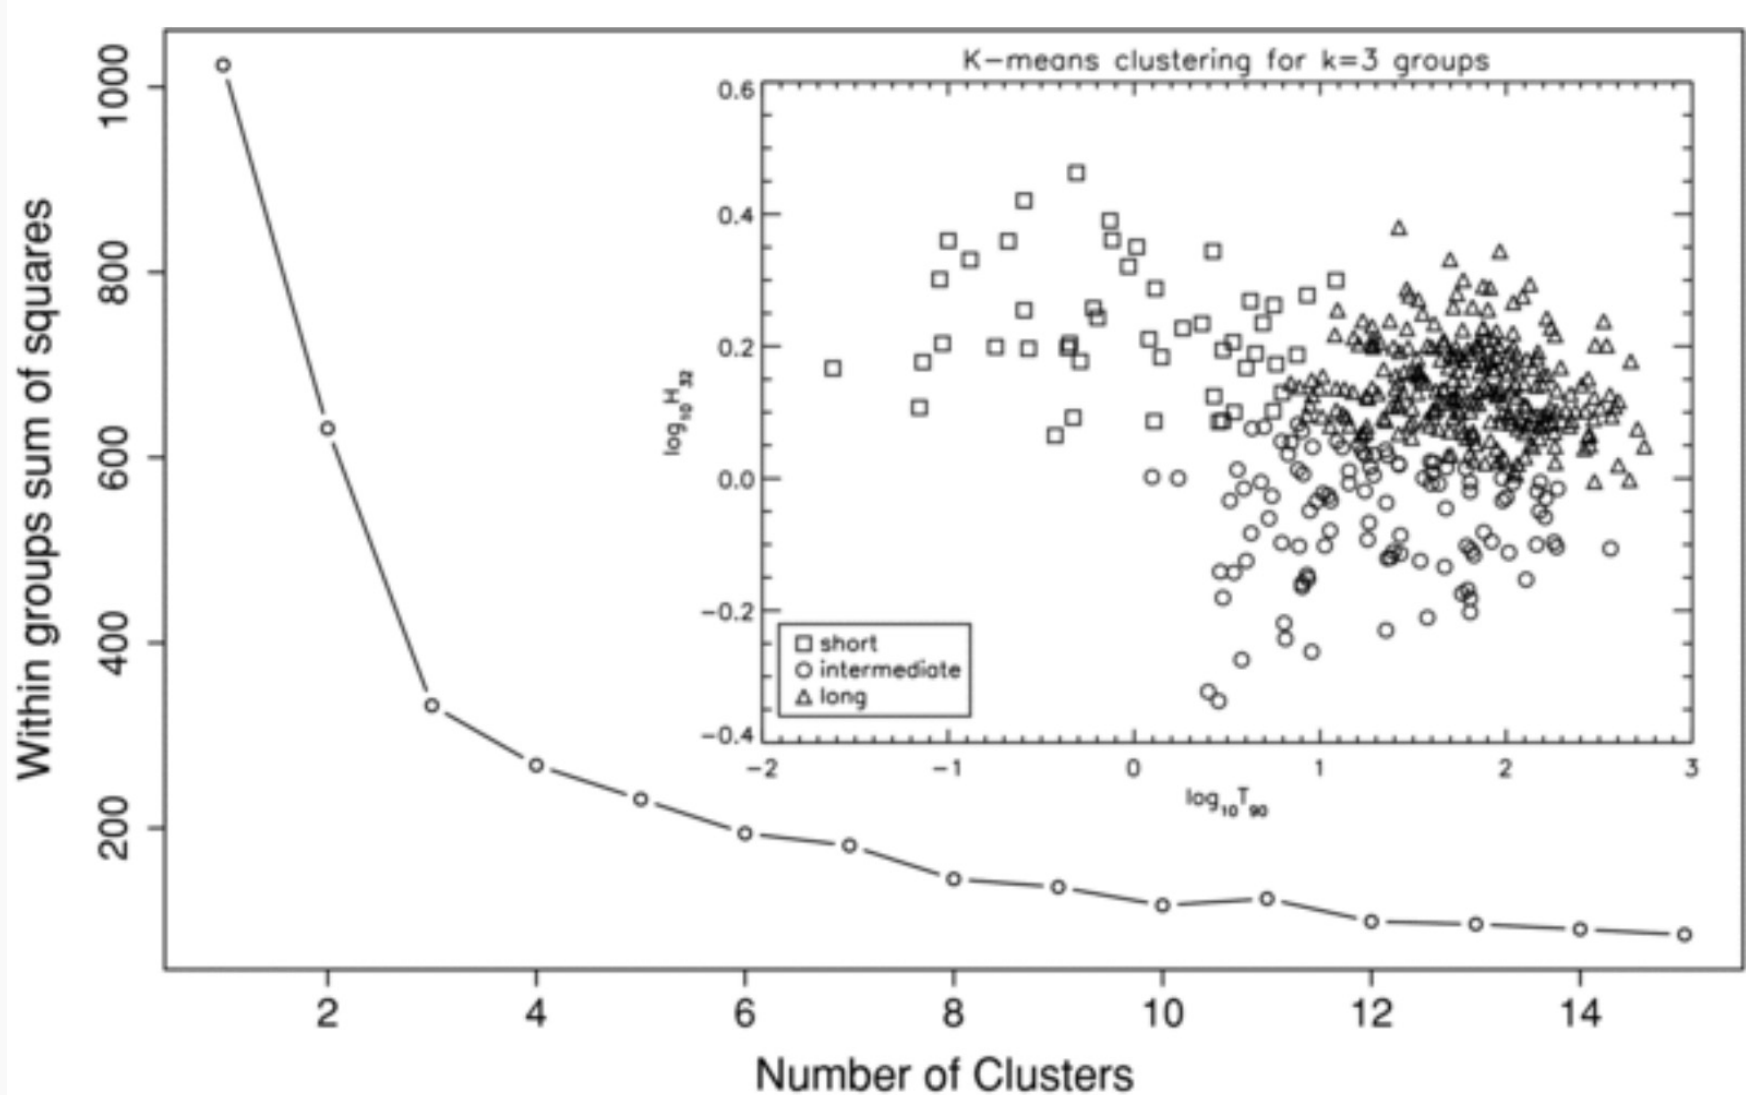
- Random Positions
- Random data points as Centers
- Random Cluster assignment to data points
  
- Start several times

# How to find K

---

- Extreme cases:
  - $K=1$
  - $K=N$
- Choose K such that increasing it does not model the data much better.

# “Knee” or “Elbow” method





# Cross Validation

---

- Use this if you want to apply your clustering solution to new unseen data
- Partition data into  $n$  folds
- Cluster on  $n-1$  folds
- Compute sum of squared distances to centroids for validation set

# Getting Rid of K

---

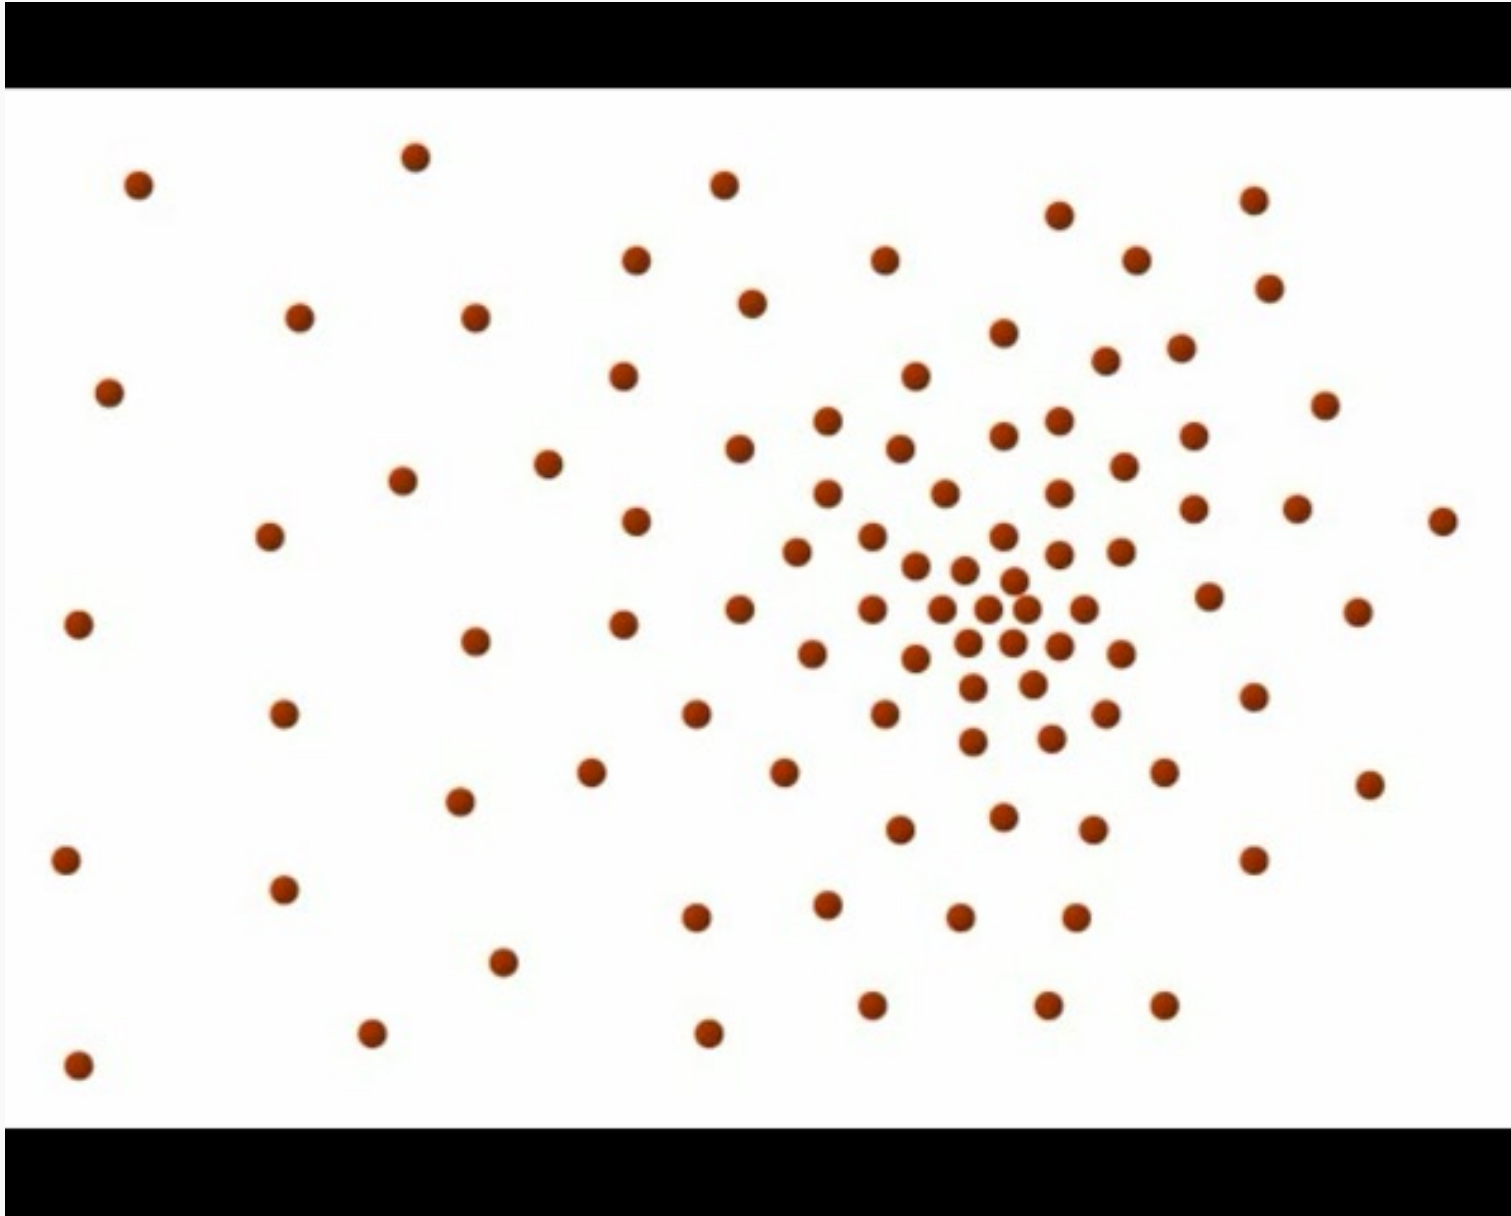
- Having to specify K is annoying
- Can we do without?

# Mean Shift

---

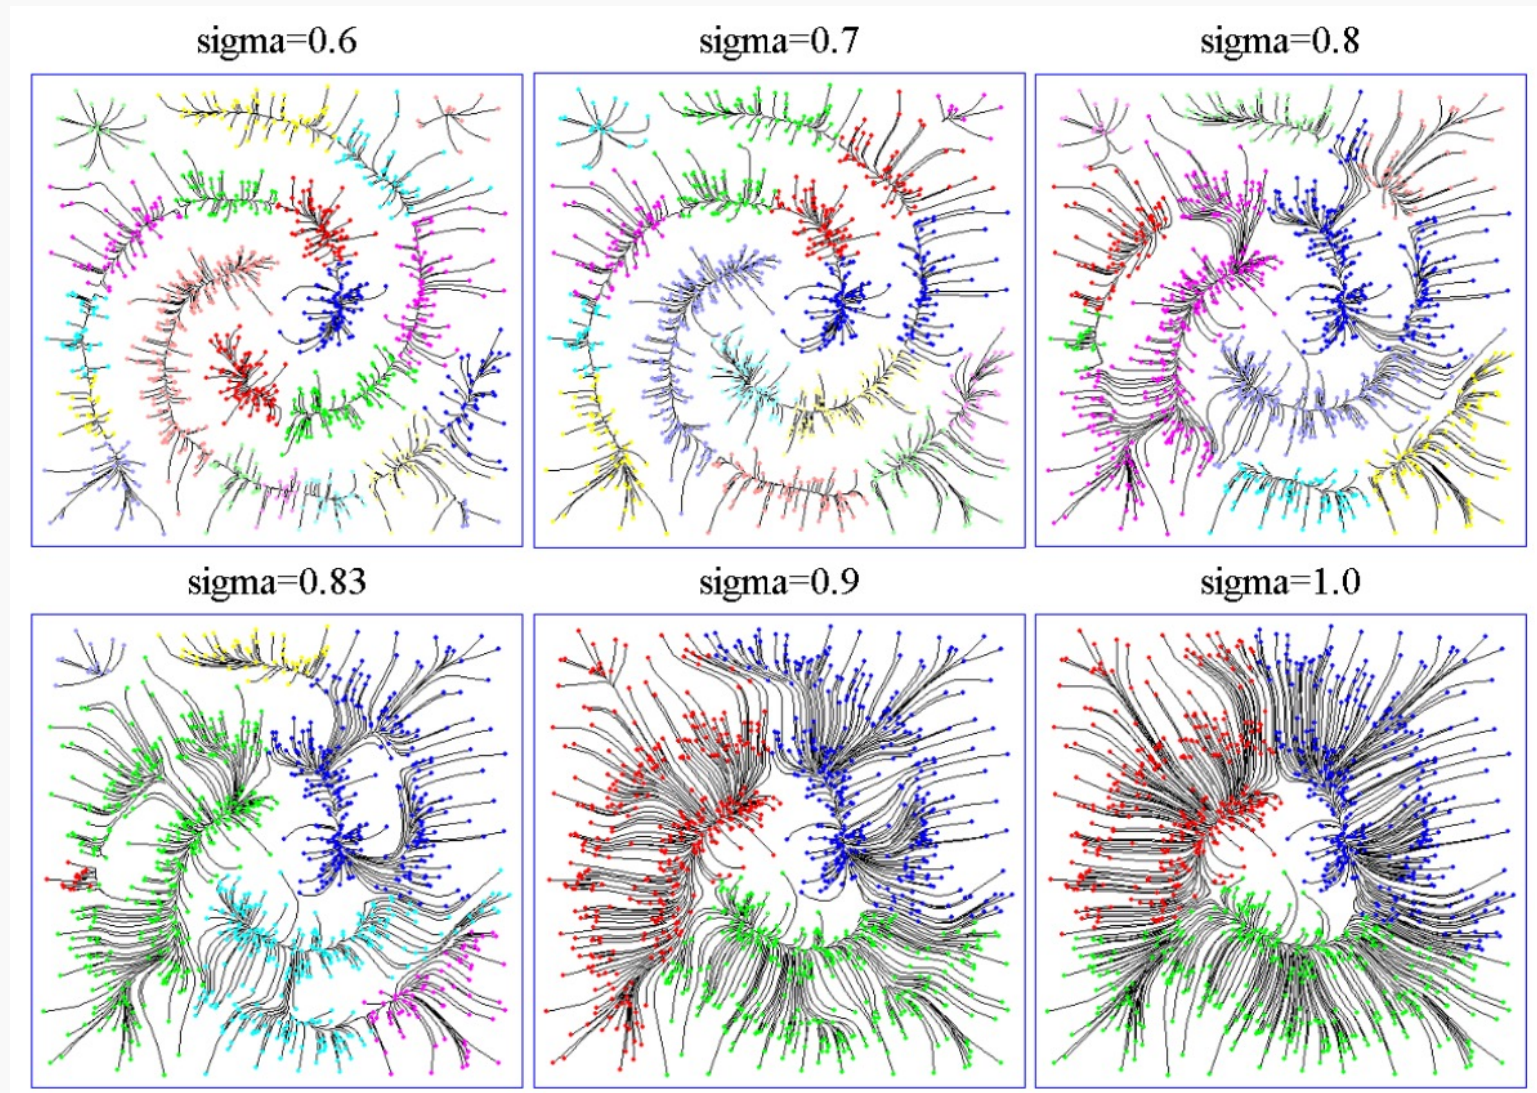
1. Put a window around each point
2. Compute mean of points in the frame.
3. Shift the window to the mean
4. Repeat until convergence

# Mean Shift



<https://www.youtube.com/watch?v=kmaQAsotT9s>

# Mean Shift



Fischer et al., "Clustering with the Connectivity Kernel", NIPS (2003)

# Mean Shift Summary

---

- Does not need to know number of clusters
  - Can handle arbitrary shaped clusters
  - Robust to initialization
  - Needs bandwidth parameter (window size)
  - Computationally expensive
- 
- Very good article:  
<http://saravananthirumuruganathan.wordpress.com/2010/04/01/introduction-to-mean-shift-algorithm/>